THE QUESTIONS YOU SHOULD BE ASKING.

Twenty questions a careful investor asks, answered with the same audited numbers as the memorandum. Nothing here is spin — where an answer has a weakness, it is stated. Several answers changed materially in this edition because the structure changed — the demand gate, the preferred return, and the removal of day-one debt are called out where they matter, not quietly absorbed.

[SECTION 01 · MARKET & RATES]

Q01 / WHAT IF B200 RENTAL RATES FALL FASTER THAN YOUR 10%/YEAR ASSUMPTION?

Then returns compress along a curve you can inspect in the model workbook — and in this edition the compression hits returns, not solvency, because there is no debt at close. The floor is printed on page 07 of the memorandum rather than hidden: at the $4.00 underwriting rate with 10%/yr decay and zero residual value, five years of distributions roughly return capital (1.05×), with investors protected in ordering by the 8% cumulative preference. Historical context: H100 rates fell roughly 35% over two years while early operators still cleared strong margins — entry price matters more than the decay slope, and all-in build cost here is $72,424 per installed GPU.

Q02 / WHY WOULD A TENANT PAY $4.50/HR WHEN AGGREGATORS LIST B200 AT $3.50?

Different products. Sub-$4 listings are shared-infrastructure, on-demand capacity — preemptible in practice, variable in performance, no dedicated fabric, and sold by hosts whose reliability varies widely. Reserved capacity here is dedicated B200s on a private NDR InfiniBand fabric, guaranteed 24/7 for a year, with a named engineer. The comparable products — reserved dedicated capacity at Lambda or CoreWeave-class providers — were verified in July 2026 at $4.50–$5.50. We are at the low end of the right comparison, not the high end of the wrong one.

Q03 / WHAT HAPPENS WHEN B300 AND NEXT-GEN CHIPS SHIP?

B200 shifts down-market and keeps earning — the same pattern as every NVIDIA generation. V100s from 2017 still rent today. The model already prices this: the revenue line declines 10% every year for five years. Note this version is *more* conservative than the earlier draft — it assumes **zero** residual hardware value at Year 5, where the previous memorandum assumed $450K. The upside version: Phase 1B and Phase 2 deploy then-current silicon on infrastructure this raise already paid for — power, fiber, cooling design, and channel relationships carry over at zero marginal cost.

Q04 / CAN'T A HYPERSCALER JUST CRUSH YOU ON PRICE?

They are going the other way. AWS lists B200 at $14.24/GPU-hour because its cost structure — metro land, retail power, enterprise sales — requires it. The threat is not hyperscalers; it is other small operators with cheap power, and the marketplaces that aggregate them. Our defense is that the cost stack is mostly owned: land at $0, power at ~$0.08/kWh, batteries at manufacturer cost, fiber locked for 60 months. A competitor must buy at retail what we already hold at cost.

Q05 / WHERE DOES DEMAND ACTUALLY COME FROM? BE SPECIFIC.

Five audited channels, in priority order: Vast.ai data-center certification (buyers pay ~$7.12/GPU-hr on-demand, verified July 2026); SF Compute exchange for committed blocks; Shadeform aggregator listing; GPUaaS.com wholesale broking at a verified $4.00–$4.68 buyer band; and whole-cluster master lease with FluidStack, Voltage Park, Crusoe or GMI. Buyer-side rates in every one of those channels are verified against published sources. What is **not** verified is the payout share to us — Vast's ~75–80% figure is platform guidance, not a contract, and is confirmed only at onboarding.

[SECTION 02 · DEMAND & CONTRACTS]

Q06 / SO THE 48 RESERVED GPUS ARE NOT YET CONTRACTED?

Correct — and the structure is built so that fact cannot reach the server capital. Tranche 1 ($1,001,121) funds a certified, demand-connected site and onboarding into all five channels. Tranche 2 ($3,634,000 — the servers, 78% of the raise) is not callable until signed capacity commitments cover at least 24 GPUs at ≥$4.00/GPU-hr, through channel contracts, a binding master-lease term sheet, or direct reservations. This restores the original project's discipline — no signatures, no spend — rebuilt for a channel model instead of six named tenants. If demand never signs, the maximum loss is tranche 1, and what tranche 1 bought — a powered, cooled, fiber-connected site on owned land — retains value. A warehouse of unrented GPUs would not.

Q07 / WHAT HAPPENS IF A CHANNEL OR A TENANT DEFAULTS MID-YEAR?

Diversification is the answer that replaced tenant concentration. Revenue spreads across five channels rather than six named tenants, so no single counterparty carries the year. Reserved capacity bills quarterly in advance, freed capacity relists on marketplace aggregators within days, and the 4 spare GPUs sit outside every contract as a failure buffer. Contracts will include standard acceleration and setoff clauses, final form with counsel.

Q08 / WHY WOULD ANYONE RENEW IN YEAR 2 INSTEAD OF CHASING NEWER CHIPS?

They might not — which is why Year-2 revenue is modeled 10% lower and payback is evaluated over hardware life, not one renewal cycle. What favors renewal: migration cost (moving a production workload is weeks of engineering), the dedicated-fabric performance a tenant has tuned for, and pricing room — cash operating cost is $1.22/GPU-hr, so we can meet the market well below today's rate and stay cash-positive on operations.

[SECTION 03 · OPERATIONS]

Q09 / A SERVER DIES AT 2 AM. WHAT ACTUALLY HAPPENS?

64 physical GPUs back 60 rentable — 4 spares live outside all contracts, so a single-GPU failure is a re-map, not an SLA breach. A full-server failure (8 GPUs) degrades capacity until repair; Supermicro next-business-day support is in the operating budget, and tenant SLAs will carry service credits sized to that window. Three staffed HPC engineers ($360,000 budgeted, up from two in the earlier plan) carry the pager, and AURA monitoring is a funded $40,000 line in CAPEX.

Q10 / OKLAHOMA SUMMERS HIT 105°F. CAN ONE RETROFITTED BUILDING REALLY COOL THIS?

Yes — because this edition bought the fourth unit instead of disclosing its absence. Four 15-ton CRAC units give 60 tons installed and 45 tons with one unit failed, against a 114.4 kW IT load: true N+1 above the load, with hot-aisle containment. The fourth CRAC, a UPS re-size allowance, and a licensed PE review of the full power and cooling design are funded CAPEX lines in tranche 1, and PE sign-off is a condition of calling tranche 2. The z1Power battery bank separately rides through utility sags without dropping load.

Q11 / WHAT IF THE GRID GOES DOWN — OR DOBSON'S FIBER GETS CUT?

Power: z1Power LFP plus Eaton 93PR modular UPS — re-sized for the 114.4 kW load at N+1 under a funded tranche-1 allowance and verified in the PE review — carries the IT load through outages long enough for orderly workload checkpointing; extended-outage generator backup is a later-phase add and is disclosed as such. Fiber: single-carrier is a real Year-1 limitation; mitigations are Dobson's enterprise SLA, wireless failover for management traffic, and a second carrier in a later phase. Tenant SLAs will exclude carrier-side outages, as is standard.

Q12 / THREE ENGINEERS FOR A $4,504,711 CLUSTER — IS THAT ENOUGH?

It is lean but defensible, and it moved deliberately. The earlier plan budgeted two engineers for a 30-GPU fleet; doubling the build to 64 GPUs across 8 servers raised it to three at $120K each — $360,000, at genuine Oklahoma market rates rather than an optimistic number. Sixty rentable GPUs is still a small fleet by HPC standards. The founder adds operating capacity at no salary load, and monitoring and orchestration software is funded separately in CAPEX.

[SECTION 04 · STRUCTURE & YOUR MONEY]

Q13 / WHAT EXACTLY AM I BUYING — AND WHERE DOES DEBT FIT?

An interest in an all-equity project (final legal form with counsel): $4,635,121 raised, called in two tranches. You receive an 8% cumulative preferred return on drawn capital, then 80% of distributions above it; the sponsor's 20% promote begins only after your pref is current. There is no debt at close. The previous edition assumed $3,603,769 of senior debt with no term sheet — financing risk stacked on execution risk — and we removed it rather than defend it. Debt appears only as an optional stabilisation refinance raised against real operating history, and nothing in the plan depends on it happening.

Q14 / WHEN AND HOW DO I GET PAID?

Quarterly, from operating cash flow, in a fixed order: operating costs and reserves first, then your 8% preferred return, then the 80/20 split. First distribution follows the first full revenue quarter — roughly month 7–8 after the demand gate clears. At the underwriting rate that is about 28% of capital distributed in Year 1; at the target rate, 32%. Because there is no lender at close, nothing ranks ahead of the pref. Quarterly reporting covers utilisation, revenue by channel, cash position, and gate status.

Q15 / HOW DO I EVER EXIT AN ILLIQUID SINGLE-SITE PROJECT?

Honestly: this is a hold-to-cash-flow investment, not a flip. Realistic liquidity paths, in order: (1) the distributions themselves; (2) a sponsor or later-phase investor buyout at a modeled multiple; (3) sale of the operating business with contracts attached — cash-flowing GPU clusters with contracted capacity are currently acquisition targets for consolidators. What we will not promise is a secondary market that does not exist.

Q16 / WHAT IS THE SPONSOR PUTTING IN, AND TAKING OUT?

In: the site, three buildings, the 3 MVA transformer, 30 acres, the 60-month fiber contract, z1Power hardware at manufacturer cost, personal operating time, and a founder's salary that is not in the budget. Out: the sponsor share of distributions per the final structure, disclosed in the term sheet before you commit. No management fees, no acquisition fees, no markup on z1Power equipment. If the project takes on senior debt, expect the lender to require a personal guarantee from the sponsor — ask to see it.

Q17 / WHAT IF THE BUILD RUNS OVER THE $604,581 CONTINGENCY?

The overrun categories that kill data-center builds — land, interconnection, core-and-shell — do not exist here because the assets exist. This edition also converts the two known engineering gaps into funded lines ($113,400: fourth CRAC, UPS re-size, optics for all 64 GPUs, PE review, securities counsel) rather than carrying them as disclosures. Residual risk is HVAC and electrical labour in rural Bryan County — capped by fixed-price bids, a tranche-1 deliverable — and NVMe pricing, which rose 70–75% quarter-on-quarter in Q2 2026 and locks at order.

[SECTION 05 · THE HARD ONES]

Q18 / ISN'T THIS ENTIRELY DEPENDENT ON ONE PERSON?

Year 1, substantially yes — that is true of every founder-operated project and we will not pretend otherwise. Mitigations: three salaried engineers run daily operations independently of the sponsor; contracts, fiber, and site sit in the project entity rather than a person; key-man insurance is priced into the $40,000 insurance line; and runbook documentation is a funded deliverable of the first 90 days.

Q19 / WHY RAISE OUTSIDE MONEY AT ALL — AND WHAT DOES THE SPONSOR MAKE?

Speed and concentration: the sponsor's capital is already the site, three buildings, the transformer, the 60-month fiber contract and z1Power inventory, all contributed without fees or markup, plus an unsalaried founder. The sponsor's return is the 20% promote above your 8% preferred return — if investors do not earn, the sponsor does not either. The honest counterpoint you should weigh: a promote rewards upside, so the sponsor is structurally incentivised toward the target rate. That is exactly why every headline number in these documents is built on the $4.00 underwriting case instead, and why the gate — not optimism — controls the server capital.

Q20 / WHAT KILLS THIS DEAL? GIVE ME THE BEAR CASE STRAIGHT.

Four things, in order. **(1) The demand gate never clears** — channel demand at ≥$4.00 never signs. Loss is capped at tranche 1 ($1,001,121, about 22% of committed capital) against a certified site that retains value — this replaces the old bear case, in which $3,634,000 of silicon had already been bought. **(2) The gate clears and rates land at the floor.** Five years of distributions roughly return capital (1.05×) and little more — the pref and the gate protected you, but the return case failed. **(3) Rates collapse beyond 10%/yr mid-hold.** The operating floor is ~$1.22/GPU-hr, so survival is not the issue; returns are. **(4) A build failure beyond contingency plus the funded fixes.** What does not kill it: one channel failing, one server dying, one bad negotiation, fiber slipping a month. You are pricing (1)–(4) against a gated, pref-protected, unlevered structure. That is the actual bet — and it is a materially better one than the previous edition offered.

[VERIFY, DON'T TRUST]

Everything above is checkable: the itemized build blueprint with its audit log, the verified demand-channel workbook with per-provider sources, the signed Dobson quote, land title, the transformer nameplate, z1Power cost invoices, and the live model workbook. The data room exists so you never have to take our word for a number.

[STILL HAVE A QUESTION?]

Ask it before you commit — an investor who stress-tests this deal is doing us a favour. Questions we cannot answer well are questions we have not de-risked yet, and we would rather find them now than in Year 2. The two engineering items disclosed in the previous edition are now funded budget lines — found by exactly this kind of stress-test.
